

UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley  
Professor  
Bora Nikolić

## Caches III

# Direct Mapped Example

# Accessing data in a direct mapped cache

- Ex.: 16KB of data, direct-mapped, 4 word blocks
  - Can you work out height, width, area?

- Read 4 addresses

1. 0x00000014
2. 0x0000001C
3. 0x00000034
4. 0x00008014

- Memory values here:

Memory Address (hex)	Value of Word
00000010	a
<u>00000014</u>	b
00000018	c
<u>0000001C</u>	d
...	...
00000030	e
<u>00000034</u>	f
00000038	g
0000003C	h
...	...
00008010	i
<u>00008014</u>	j
00008018	k
0000801C	l
...	...

# Accessing data in a direct mapped cache

- 4 Addresses:
  - 0x00000014, 0x0000001C,  
0x00000034, 0x00008014
- 4 Addresses divided (for convenience) into **Tag**, **Index**, **Byte Offset** fields

00000000000000000000 0000000001 0100

00000000000000000000 0000000001 1100

00000000000000000000 0000000011 0100

000000000000000010 0000000001 0100  
**Tag**                   **Index**                   **Offset**

# Example: 16 KB Direct-Mapped Cache, 16B blocks

- Valid bit: determines whether anything is stored in that row (when computer initially powered up, all entries invalid)

Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0				
1	0				
2	0				
3	0				
4	0				
5	0				
6	0				
7	0				
...					
1022	0				
1023	0				

# 1. Read 0x000000014

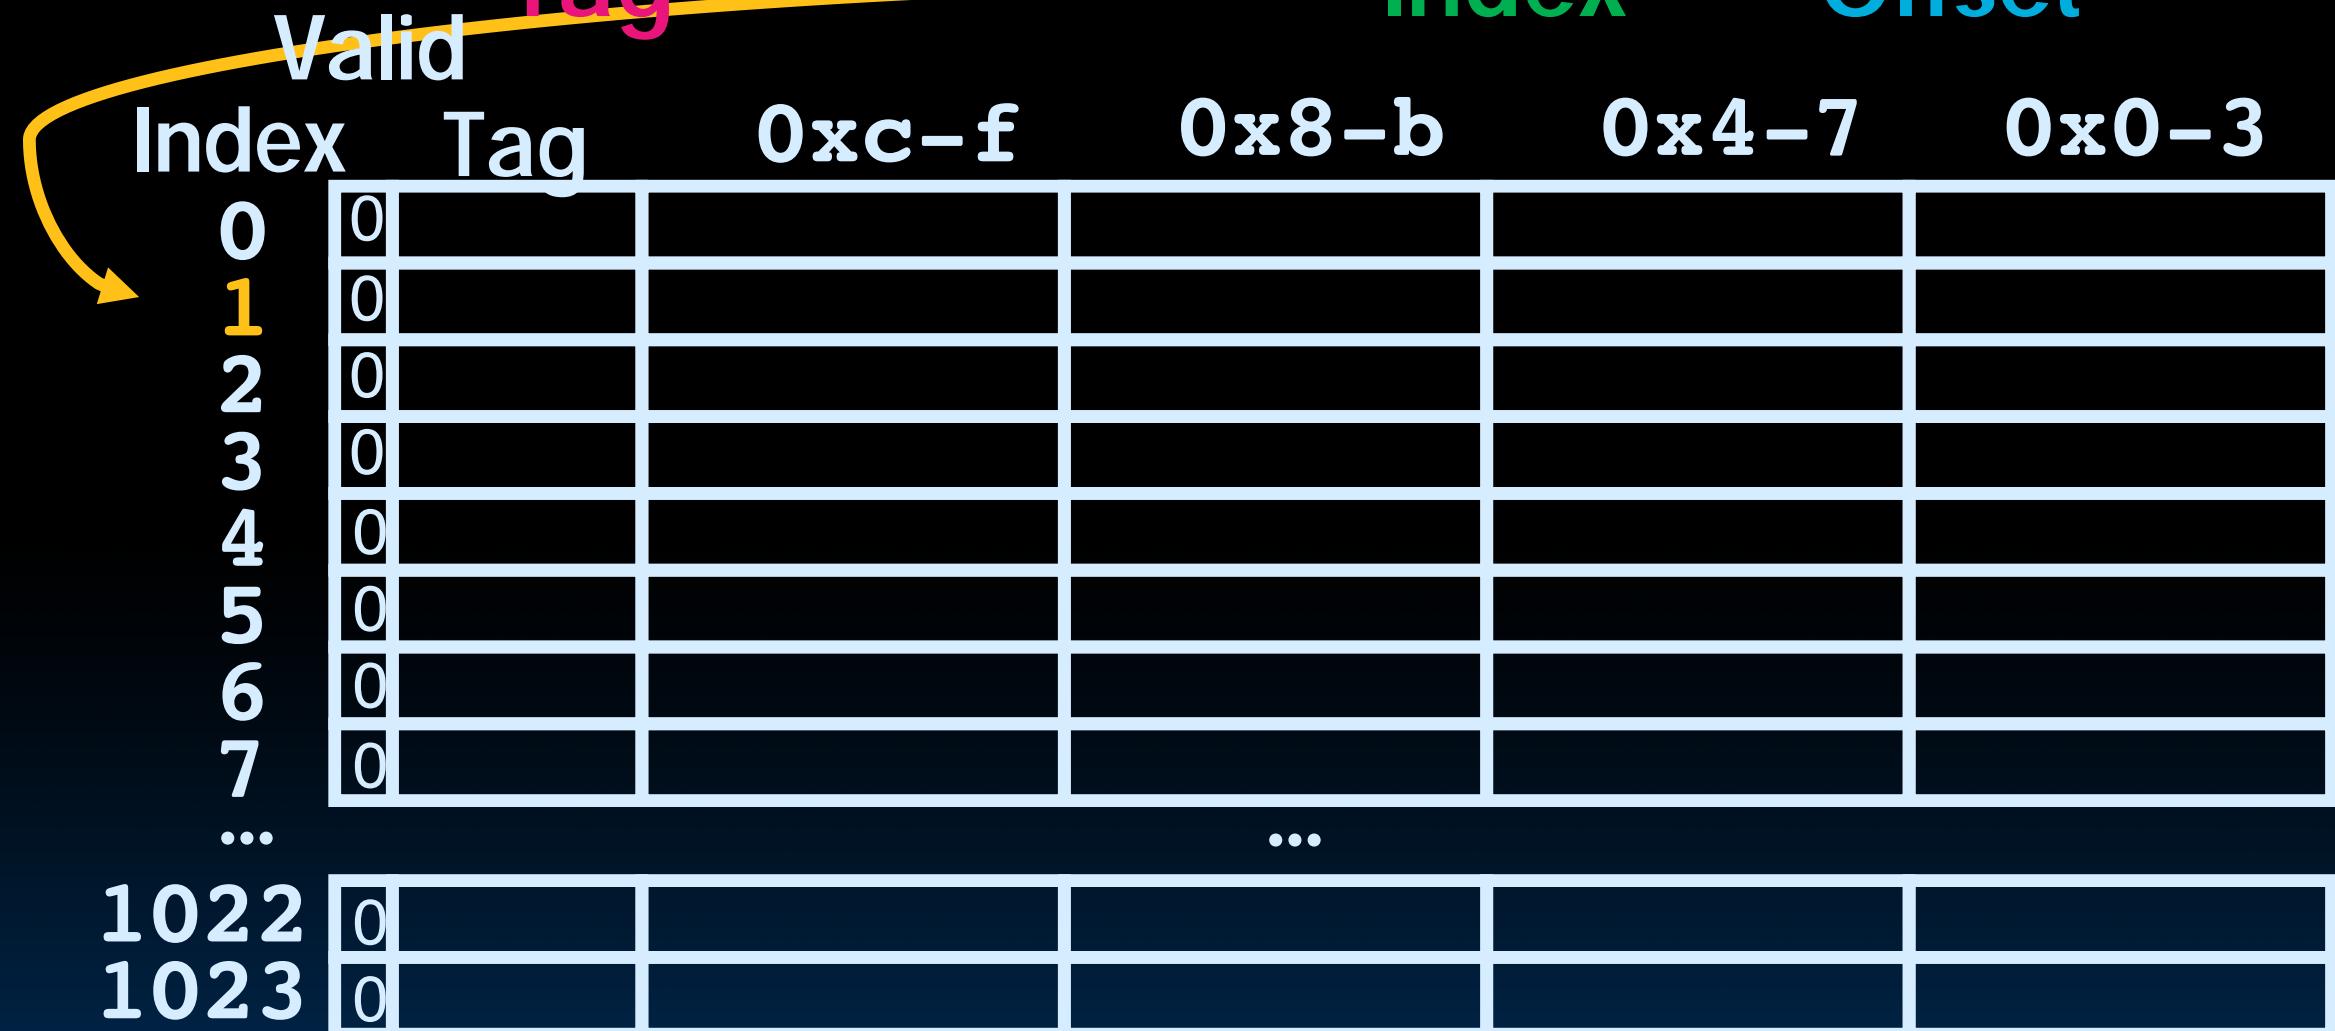
▪ 00000000000000000000000000000000 0000000001 0100  
Tag Index Offset

Valid

Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0				
1	0				
2	0				
3	0				
4	0				
5	0				
6	0				
7	0				
...					
1022	0				
1023	0				

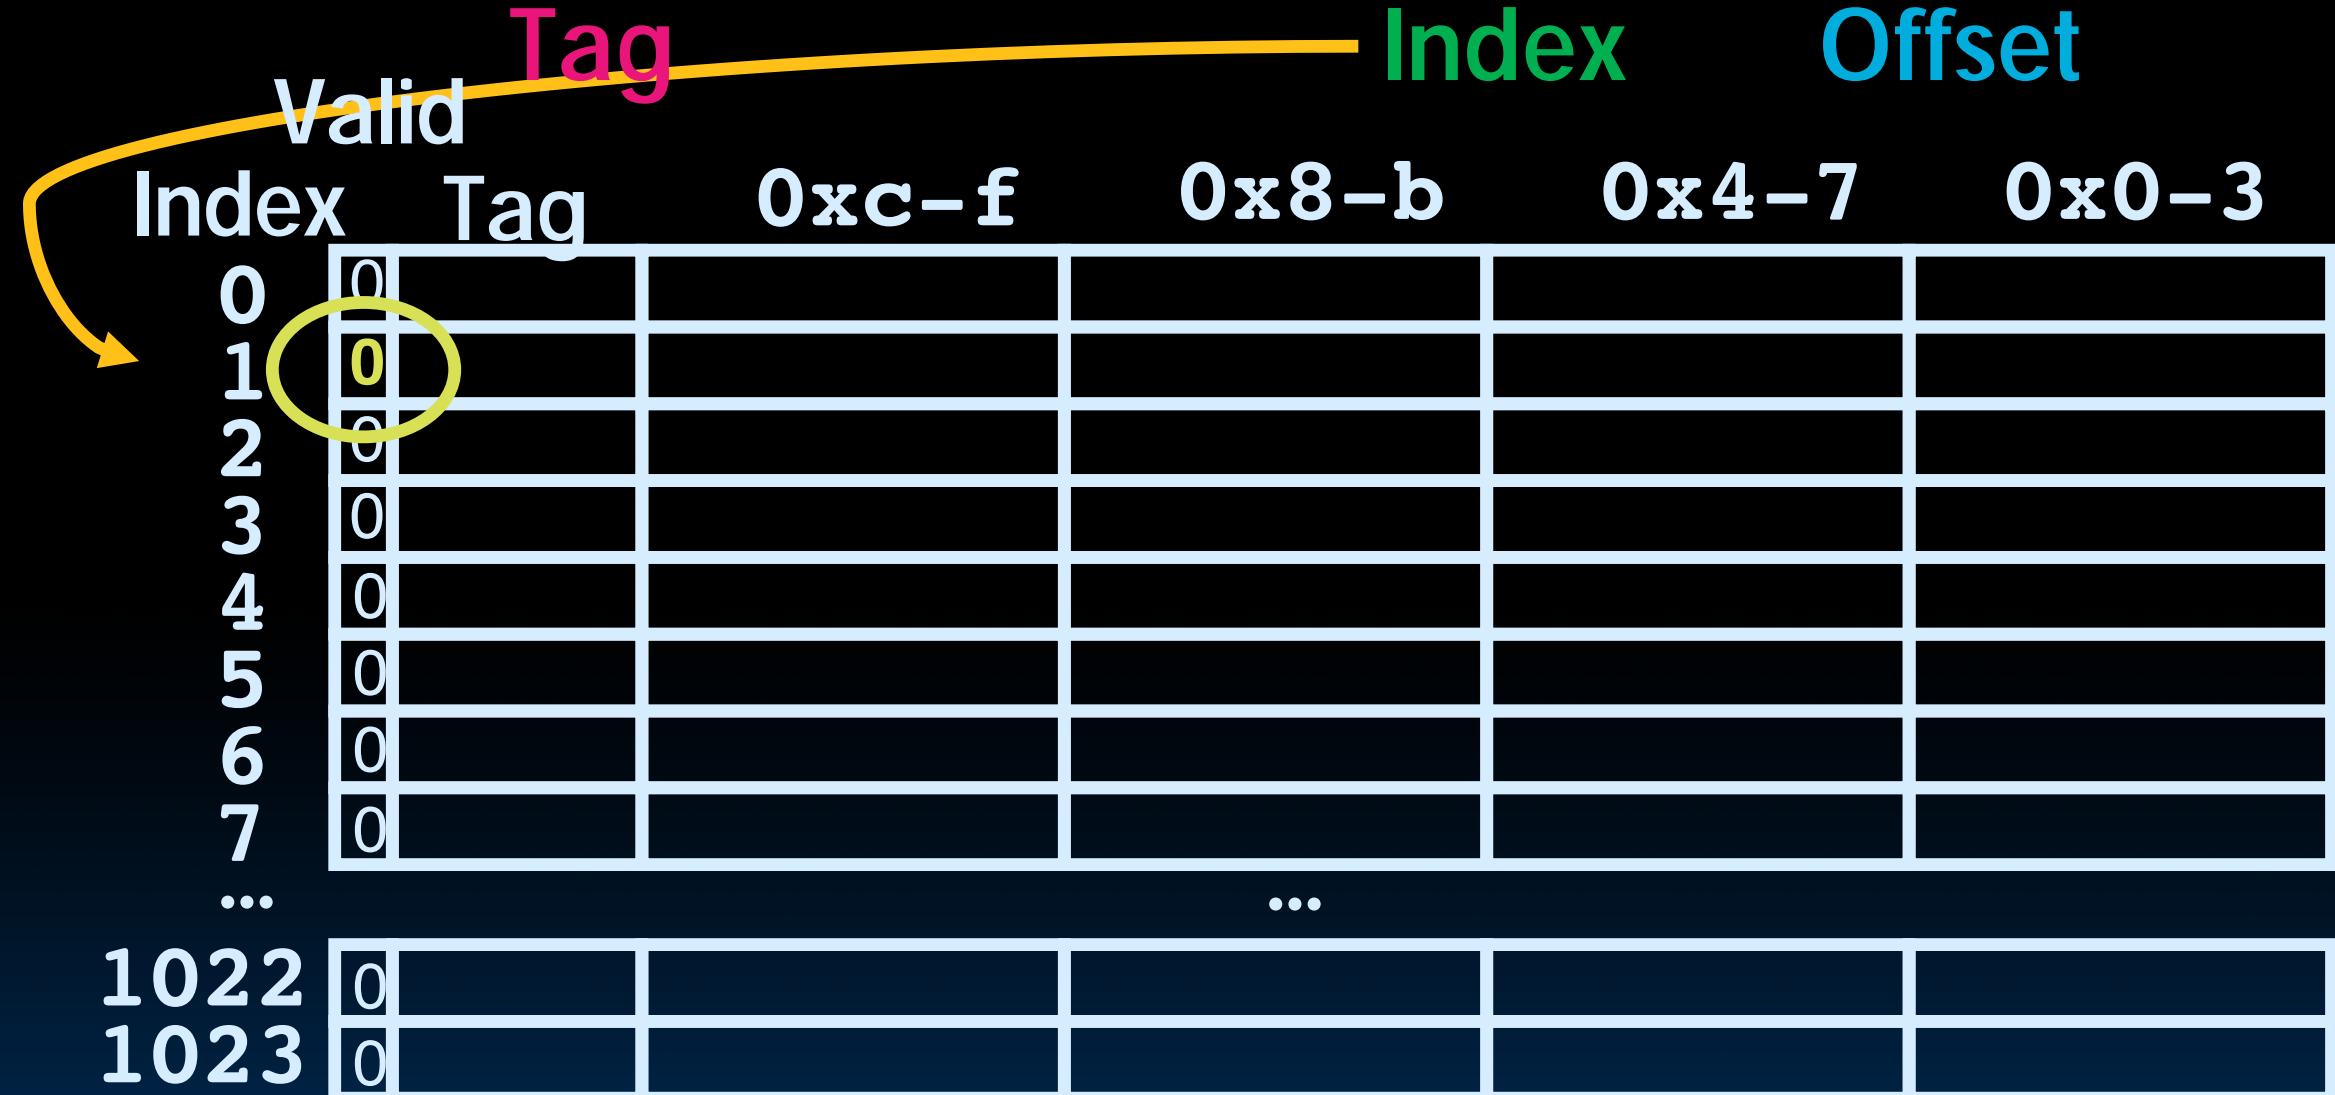
# So we read block 1 (000000000001)

▪ 00000000000000000000 0000000001 0100  
Tag Index Offset



# No valid data

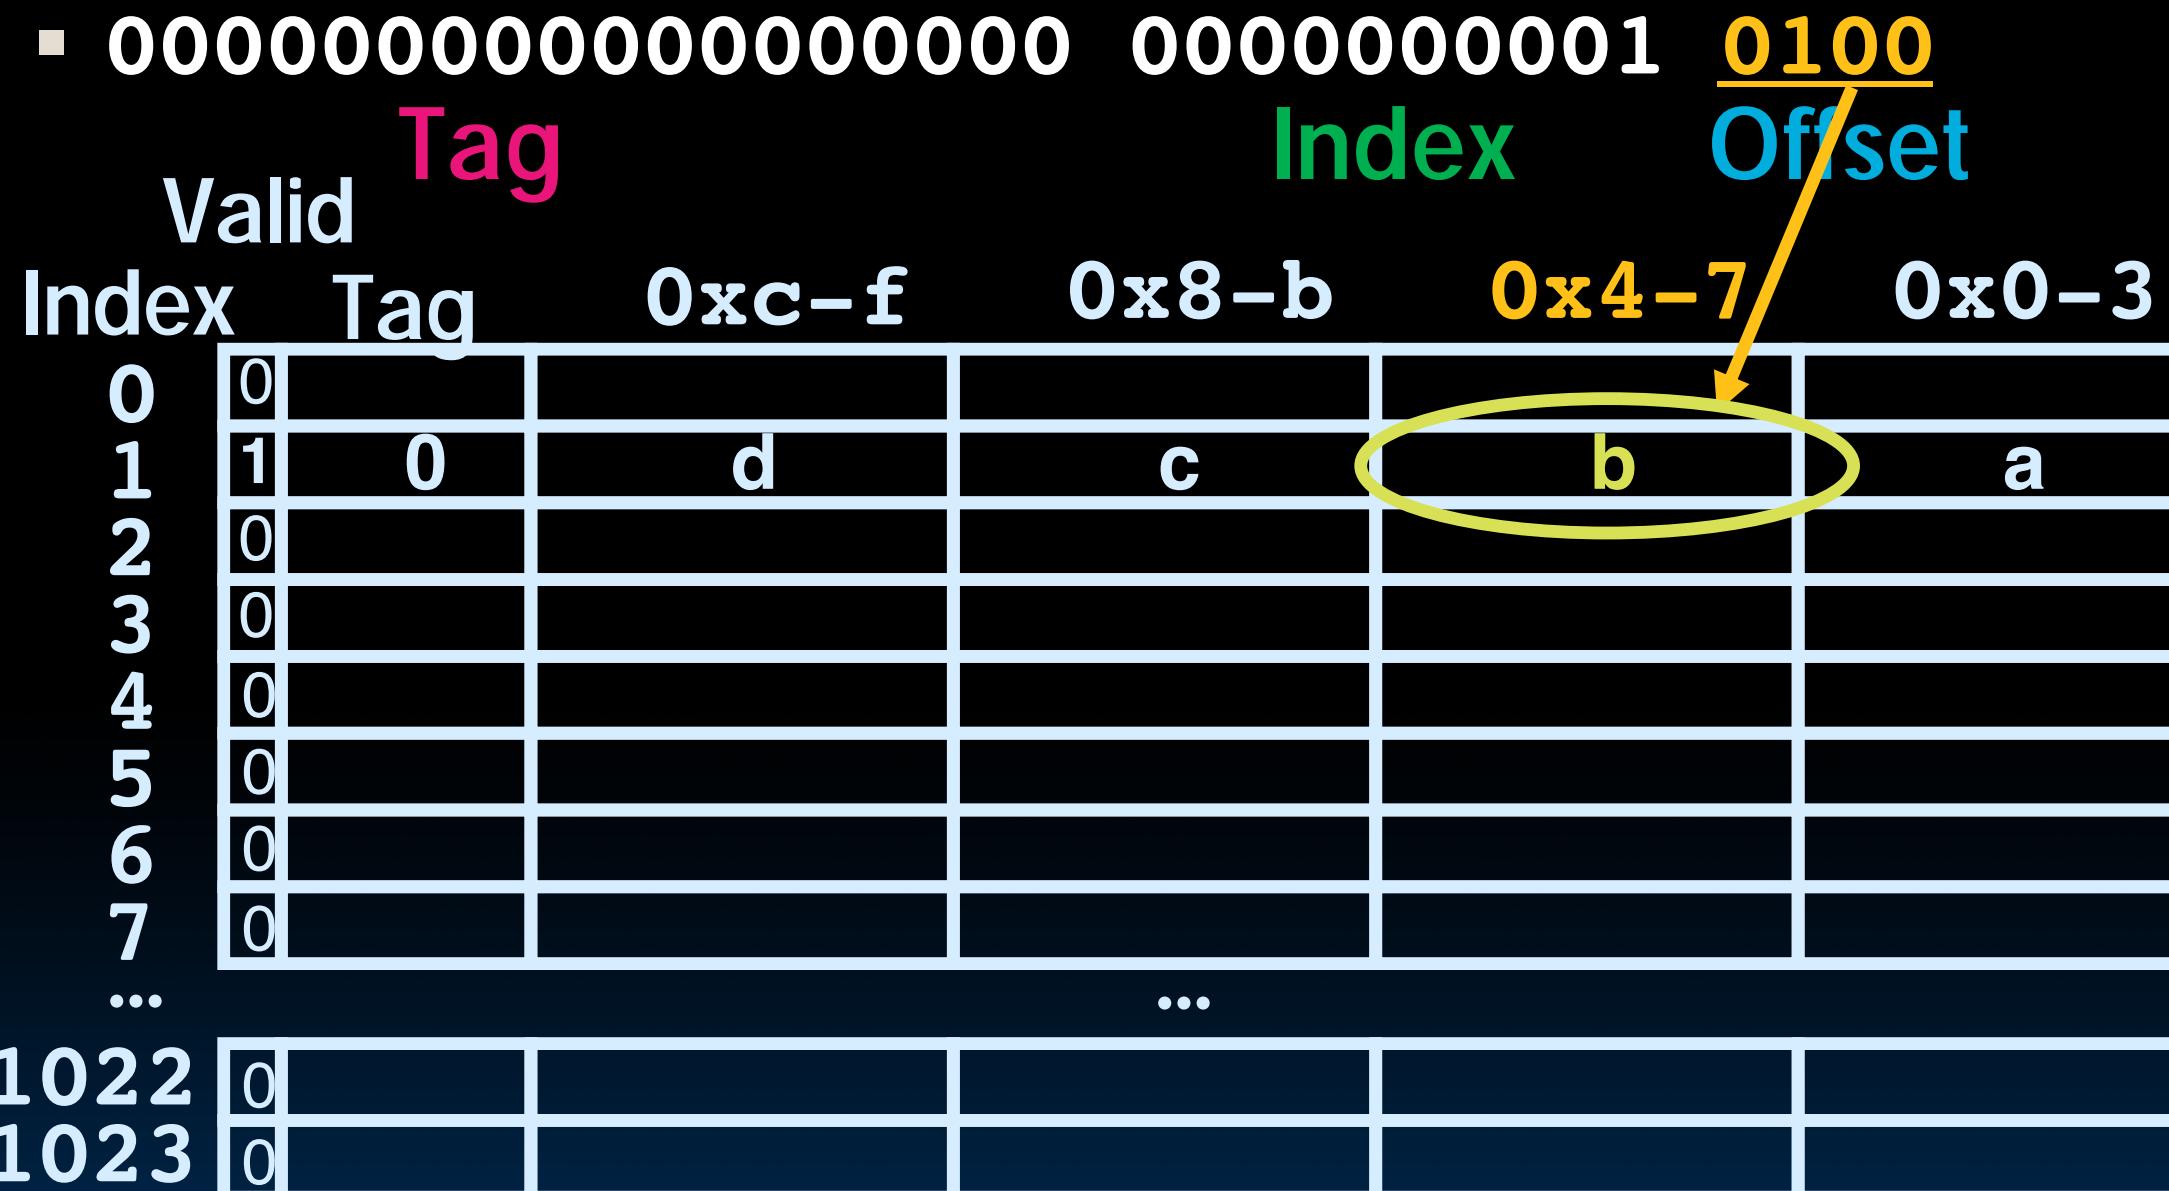
- 00000000000000000000 0000000001 0100



# So load that data into cache, setting tag, valid



# Read from cache at offset, return word **b**



## 2. Read $0x00000001C = \dots 00\ 0..01\ 1100$

- 00000000000000000000000000000000 0000000001 1100

Index	Valid	Tag	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0						
1	1	0		d	c	b	a
2	0						
3	0						
4	0						
5	0						
6	0						
7	0						
...							
1022	0						
1023	0						

# Index is Valid

- 00000000000000000000 0000000001 1100

The diagram illustrates a cache organization with 1024 rows (index 0 to 1023) and 6 columns (0xc-f, 0x8-b, 0x4-7, 0x0-3). The first row (index 0) contains all zeros. The second row (index 1) has a '1' in the 'Index' column, labeled 'Valid'. The 'Tag' column contains '0', and the other four columns contain letters 'd', 'c', 'b', and 'a' respectively. Ellipses indicate intermediate rows and the end of the cache.

	Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0	0				
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...			...			
1022	0					
1023	0					

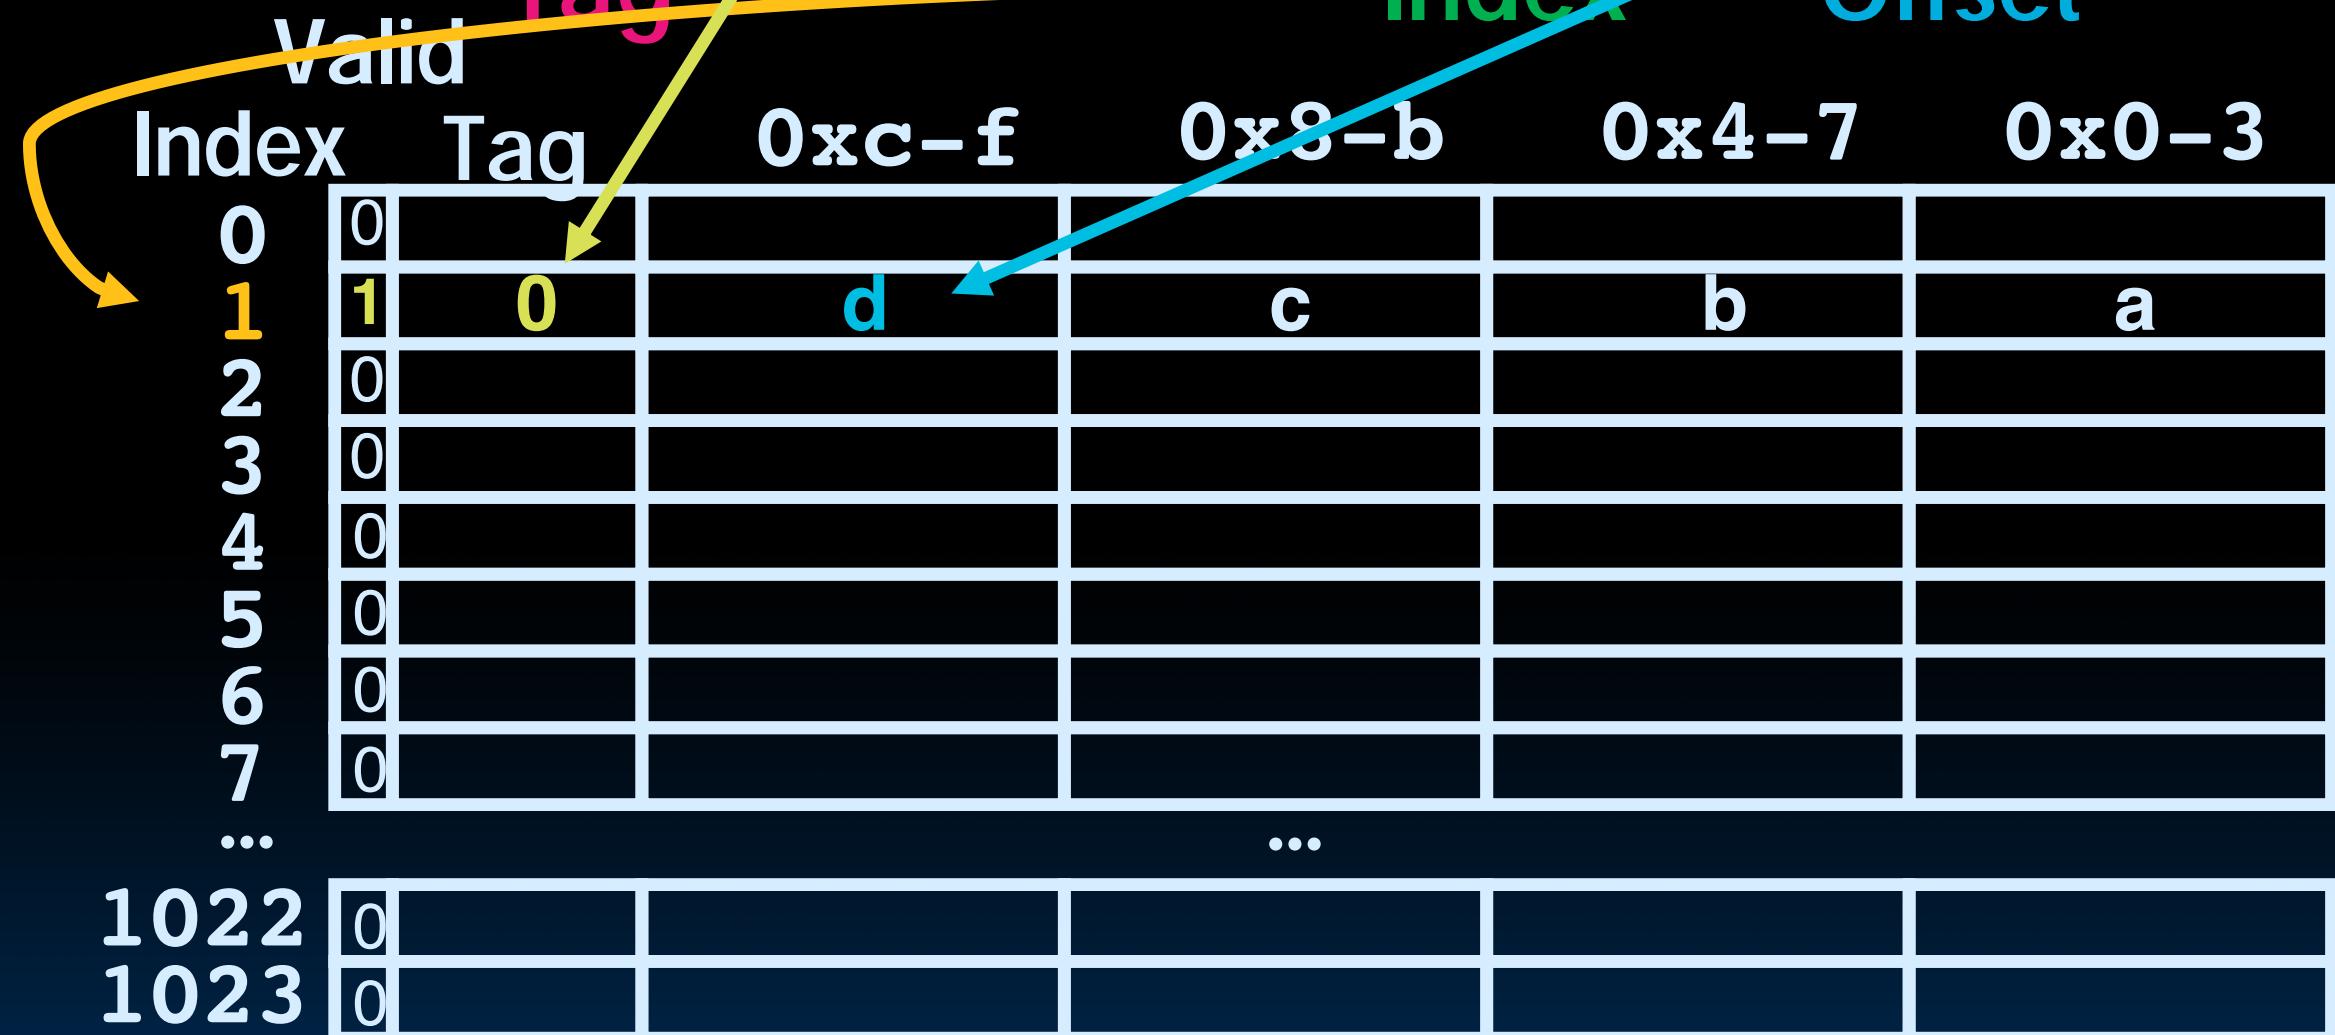
# Index is Valid, Tag Matches

- 00000000000000000000 0000000001 1100  
Tag Index Offset



# Index is Valid, Tag Matches, return d

- 00000000000000000000 0000000001 1100  
Tag Index Offset



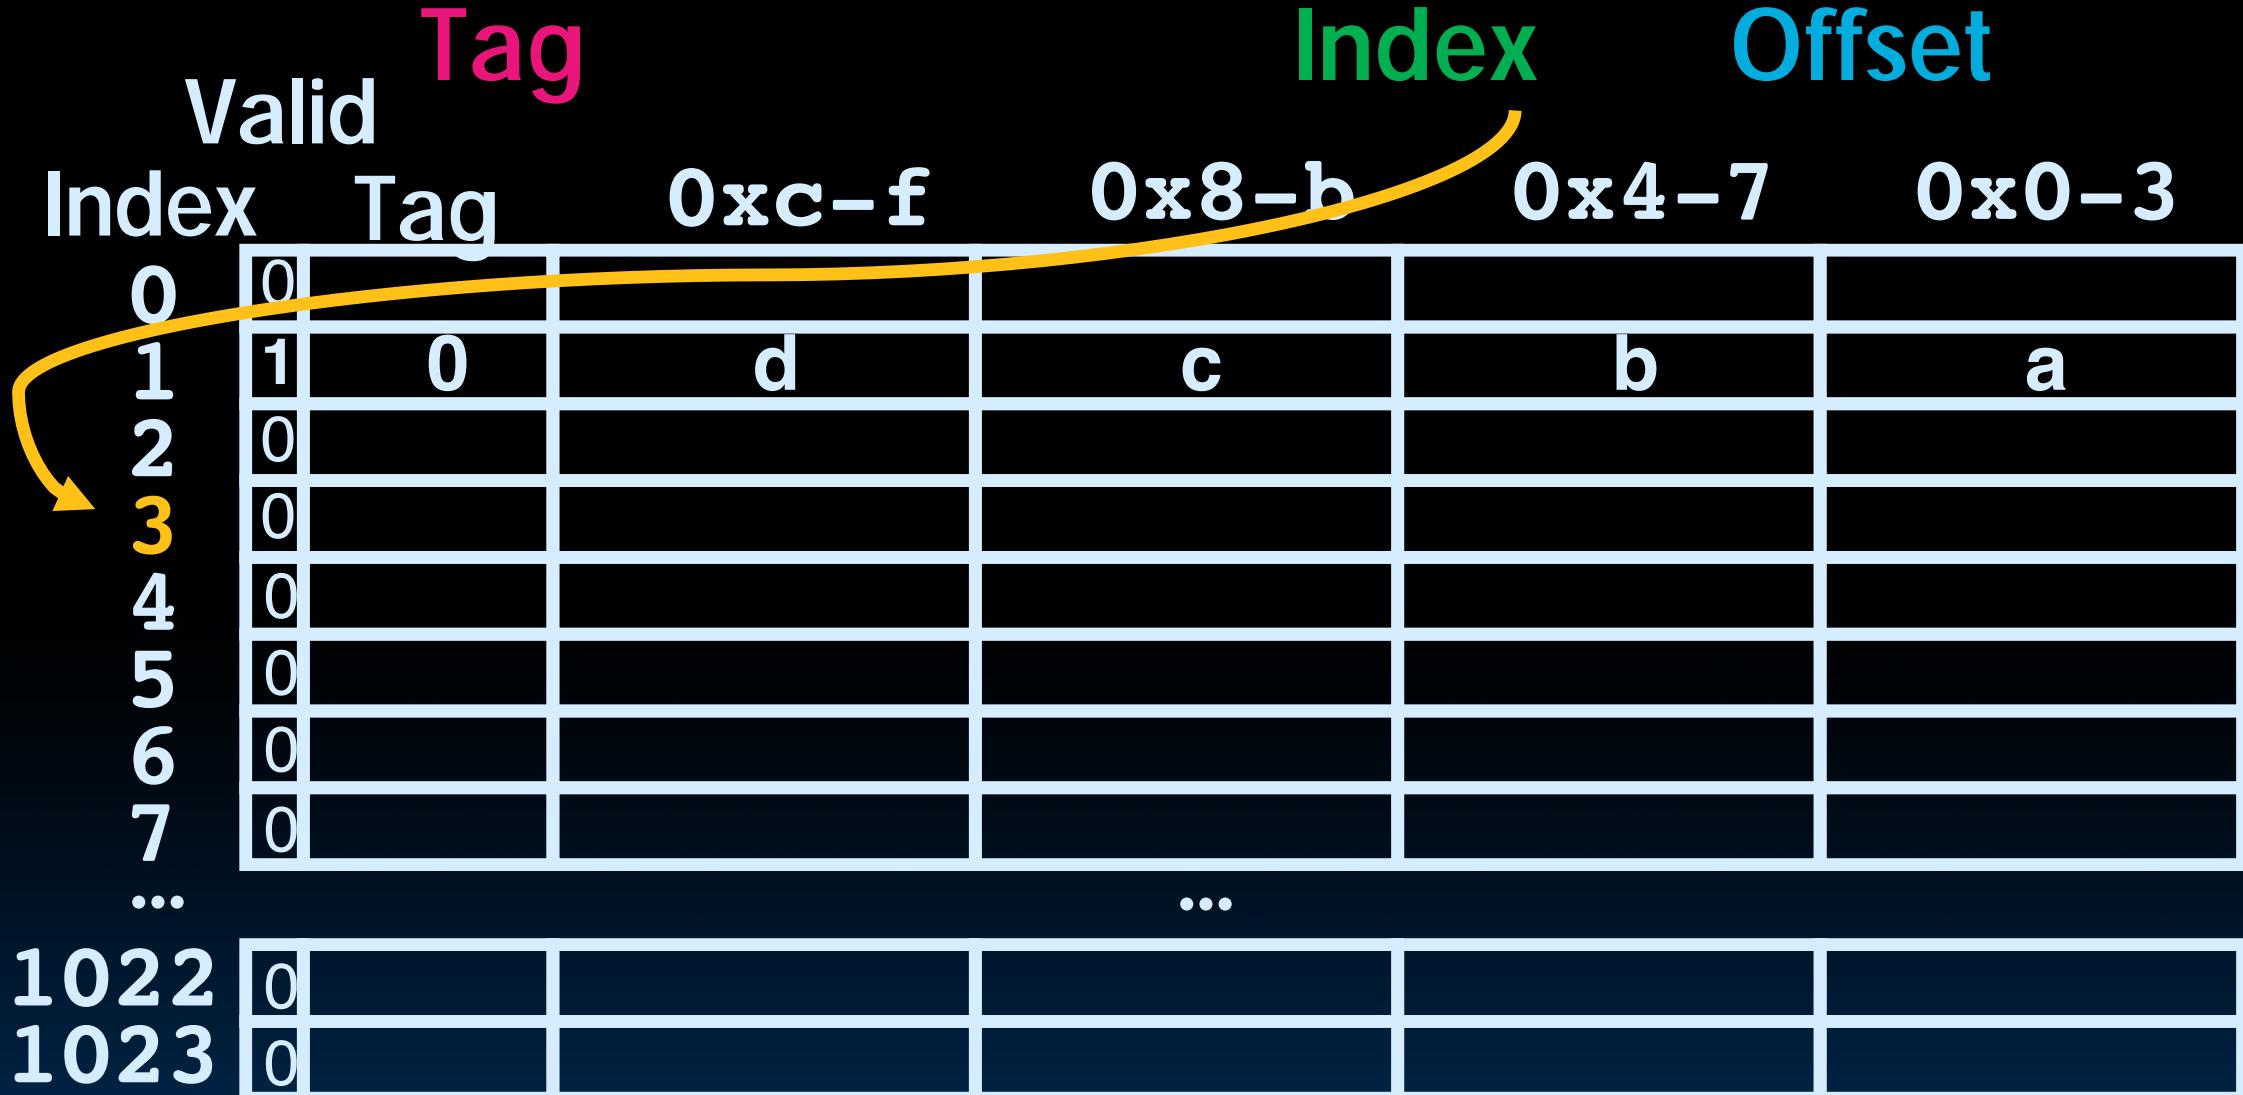
### 3. Read $0x000000034 = \dots 00\ 0..011\ 0100$

- 00000000000000000000 0000000011 0100

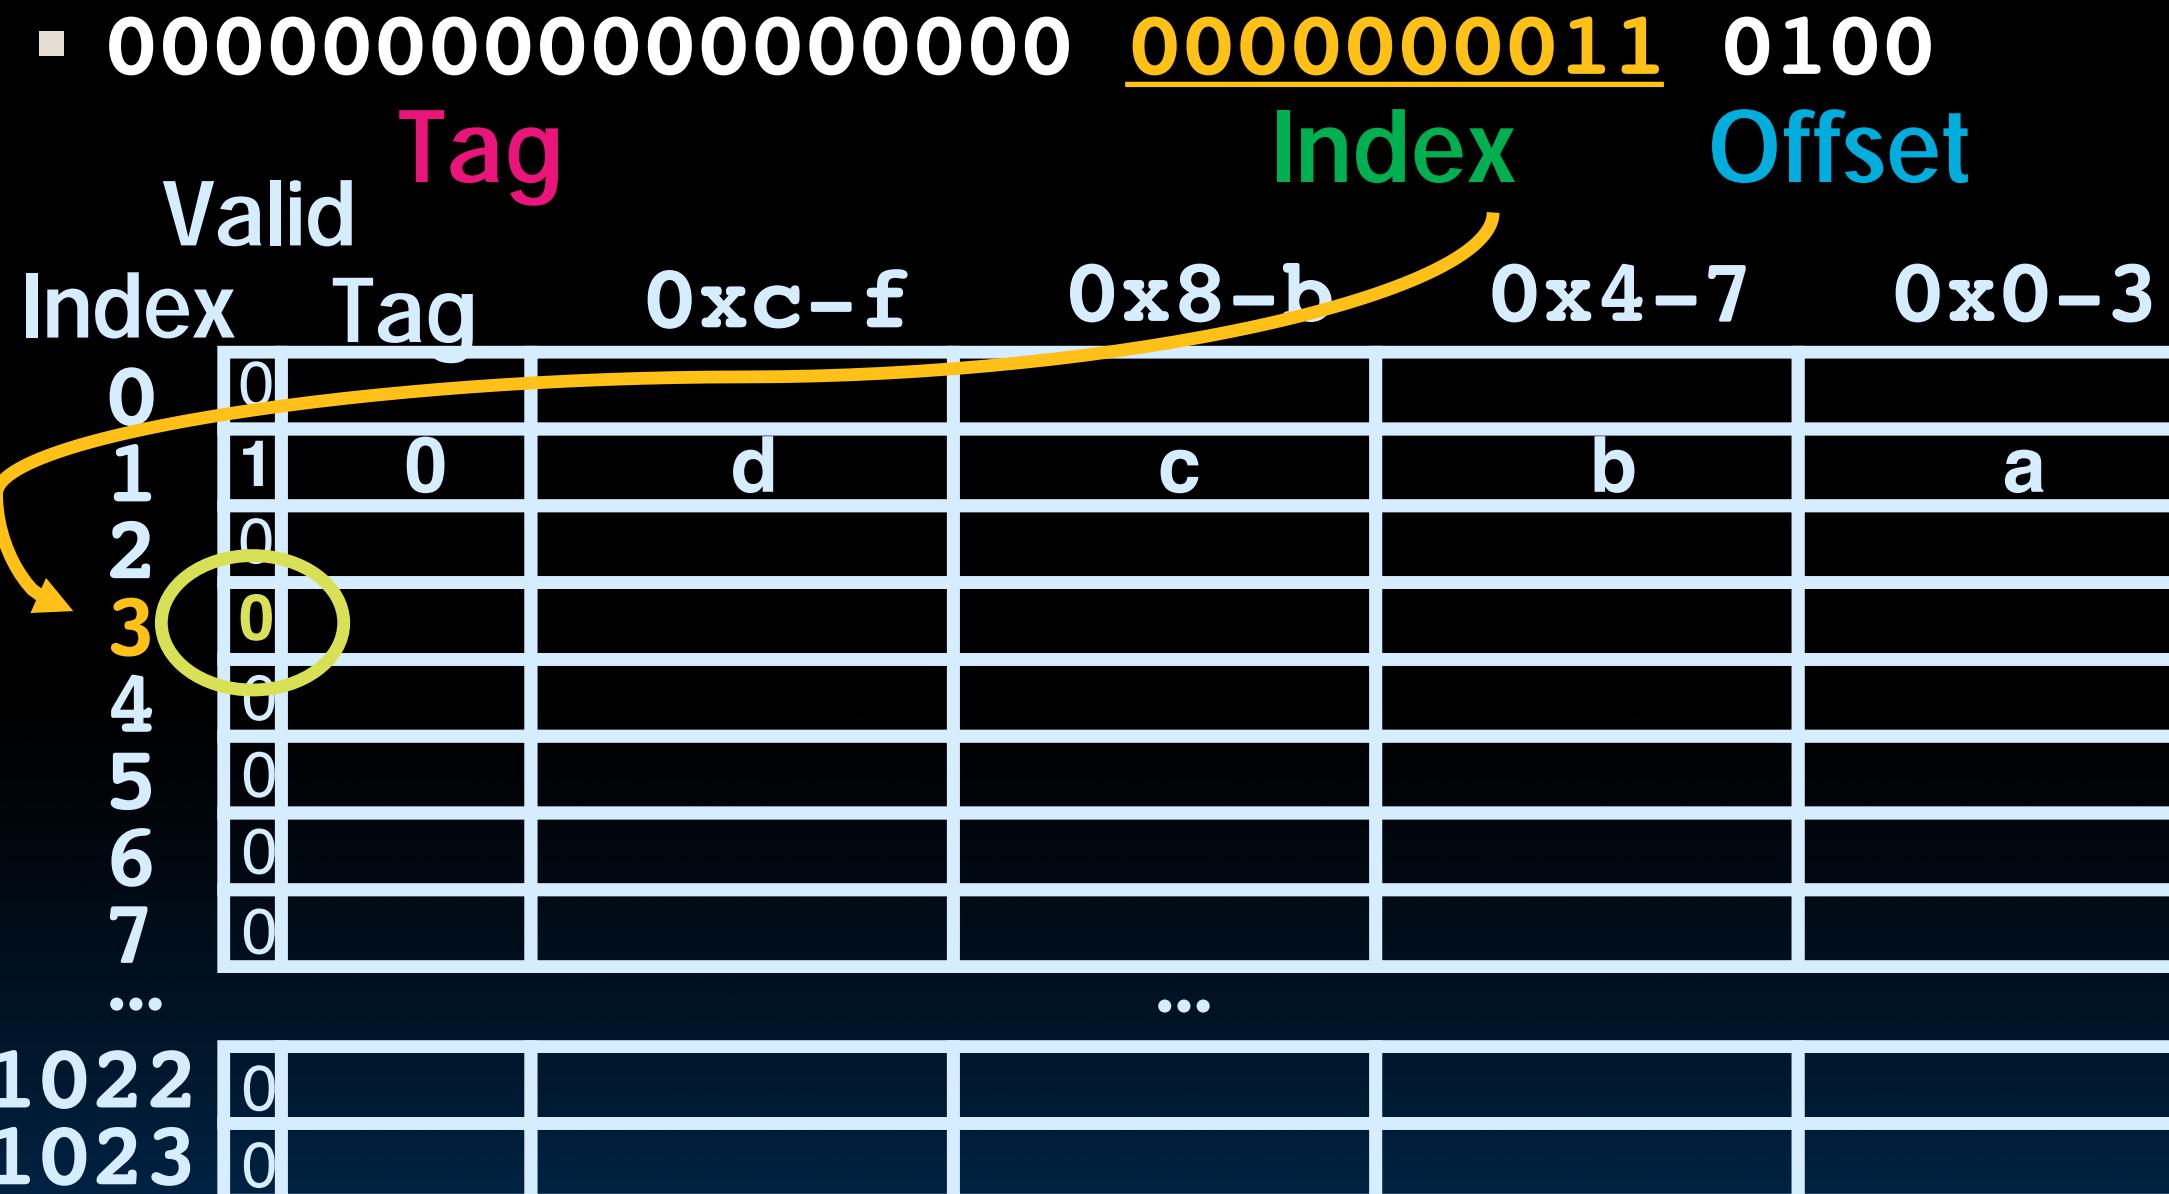
Valid	Tag	Index	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	0	d		c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...						
1022	0					
1023	0					

# So read block 3

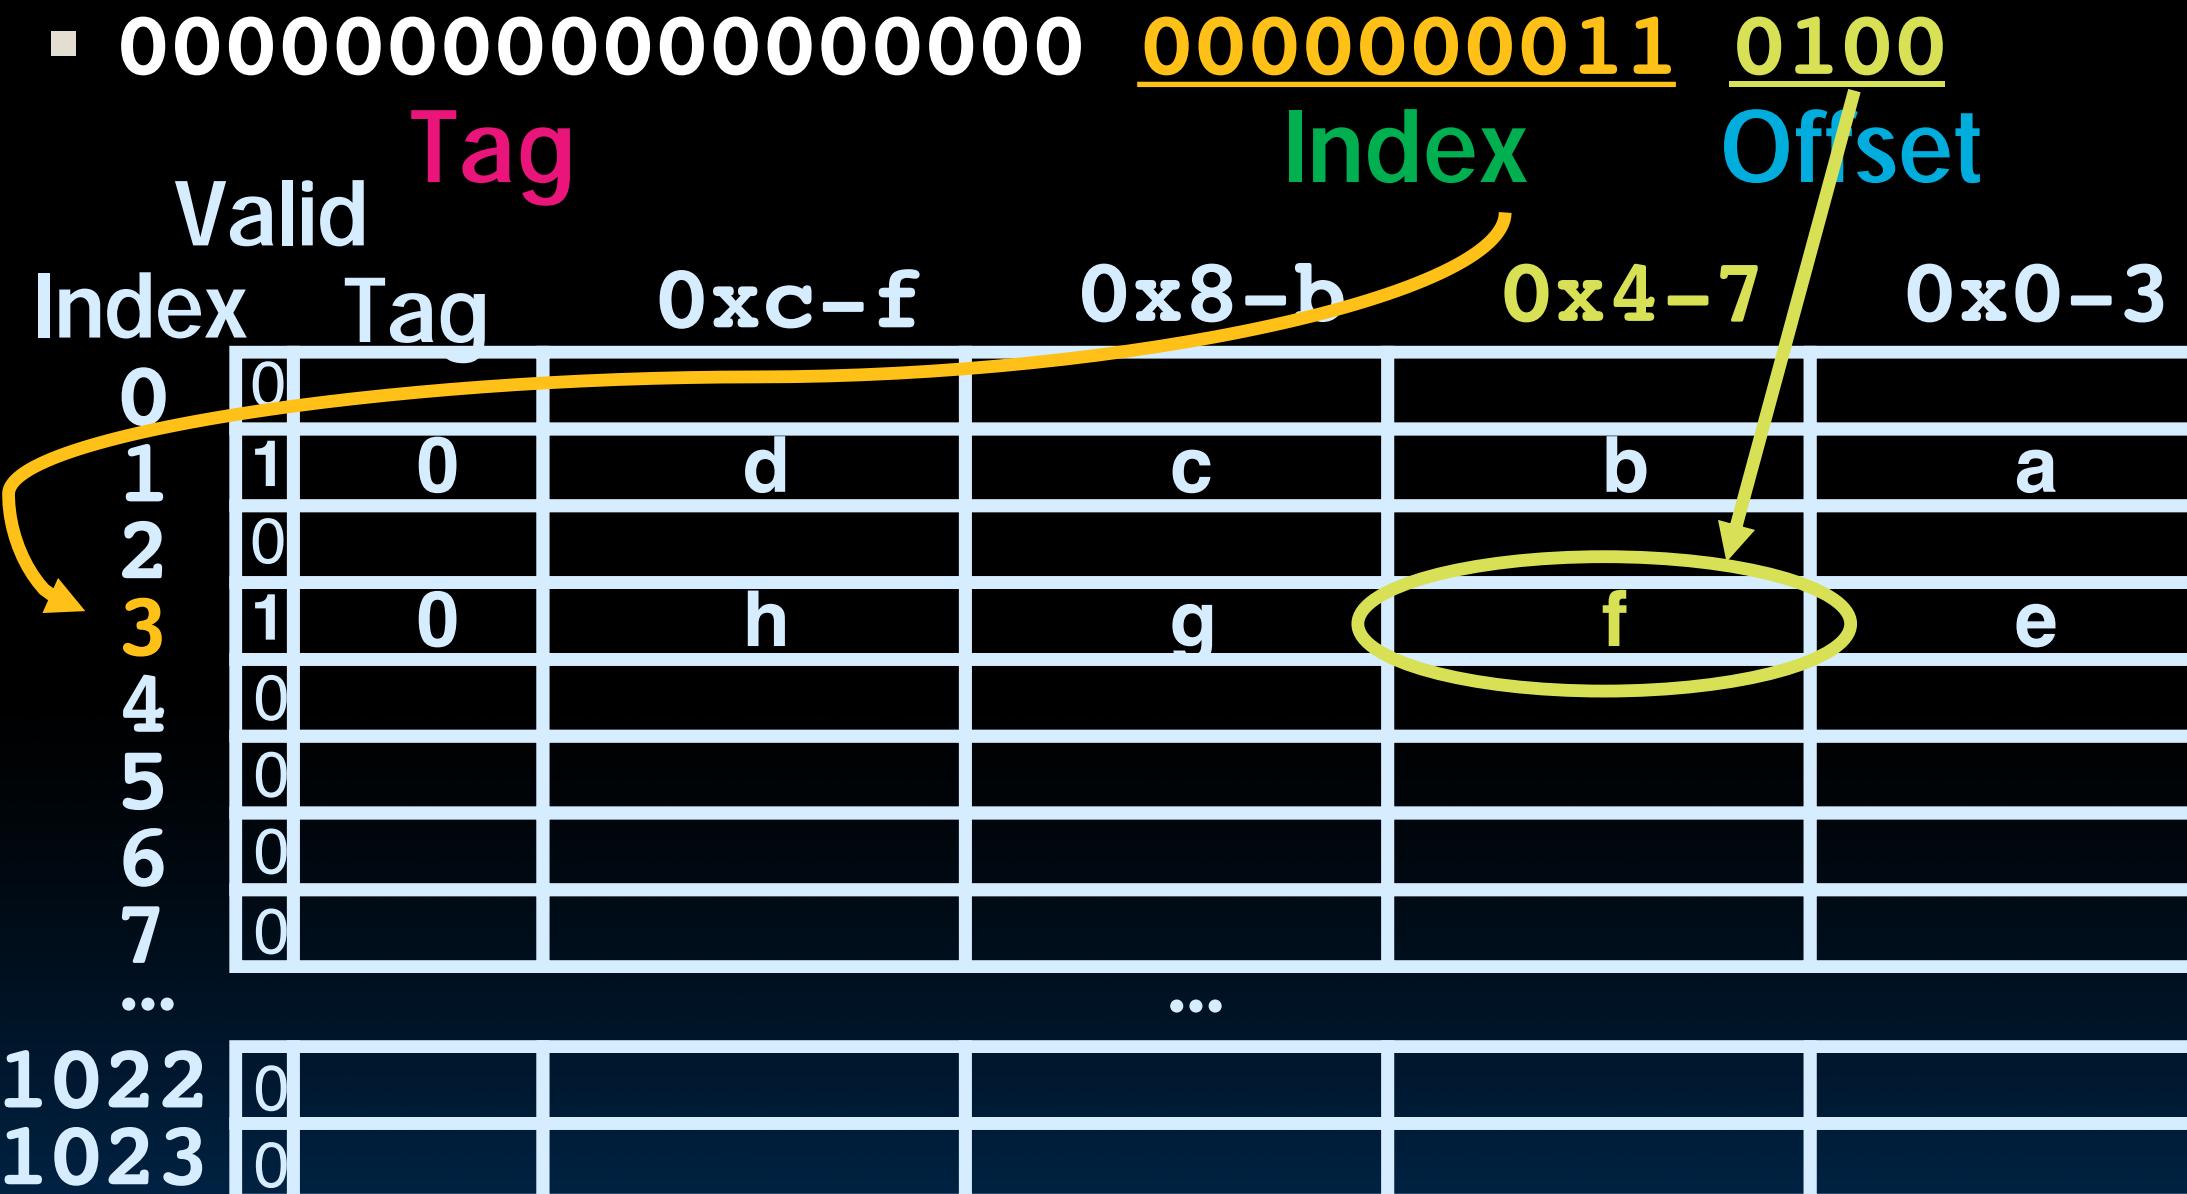
- 00000000000000000000 0000000011 0100  
Tag Index Offset



# No valid data



# Load that cache block, return word **f**



# 4. Read $0x00008014 = \dots 10\ 0..001\ 0100$

- 000000000000000010 0000000001 0100

Valid	Tag	Index	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	0	d		c	b	a
2	0					
3	0	h		g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

# So read Cache Block 1, Data is Valid

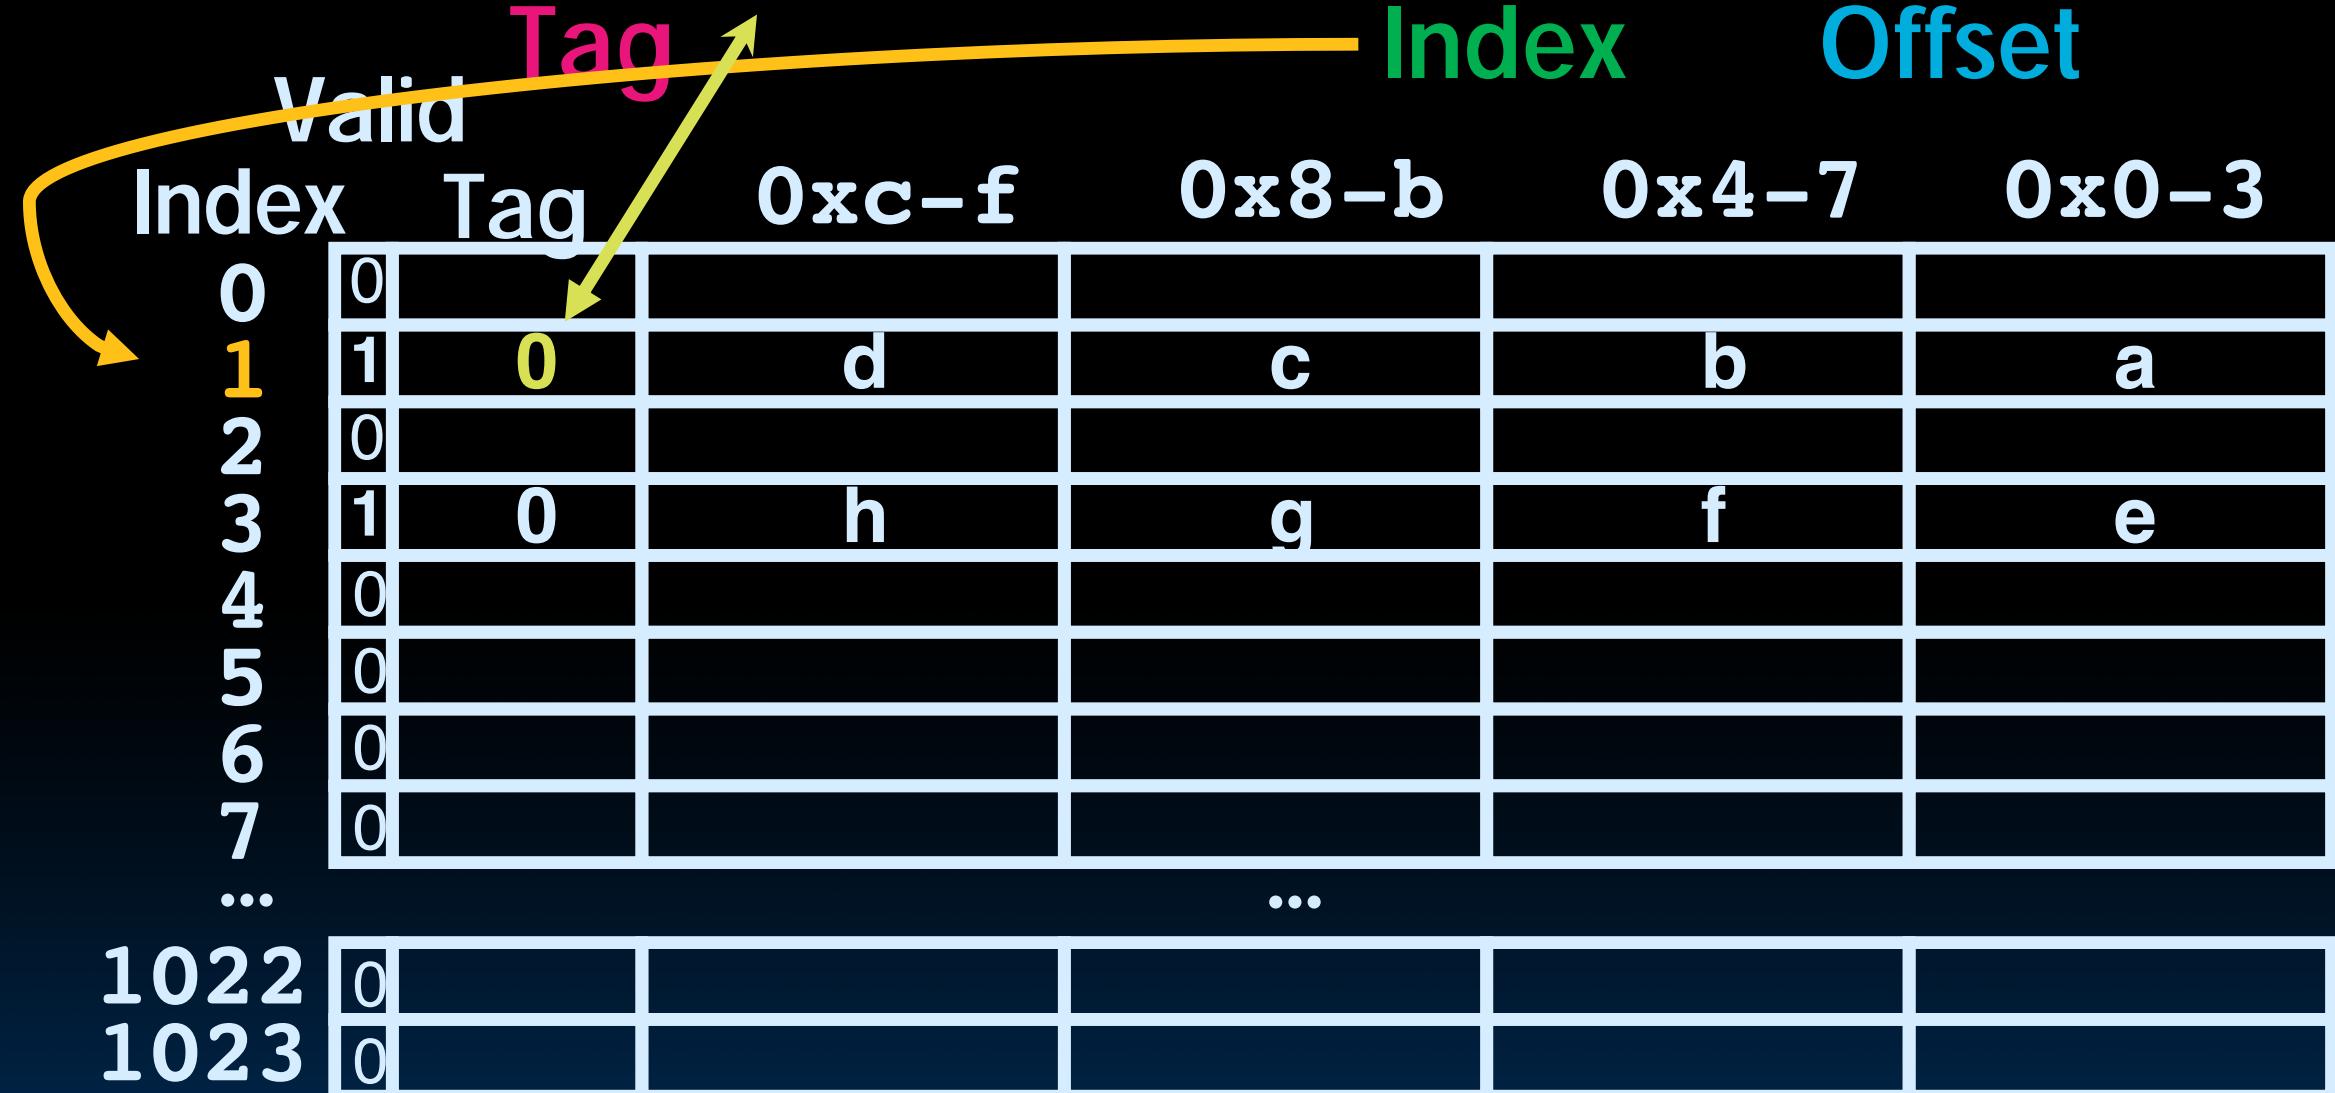
- 0000000000000000000010 0000000001 0100

The diagram illustrates a cache organization with 1024 rows (Index 0 to 1023) and 6 columns. The columns are labeled: Index, Tag, 0xc-f, 0x8-b, 0x4-7, and 0x0-3. A yellow arrow points from the search value '0000000000000000000010' to the 'Index' column. The 'Index' column contains binary values from 0 to 1023. The 'Tag' column contains binary values 0 or 1. The data columns (0xc-f, 0x8-b, 0x4-7, 0x0-3) contain characters representing data blocks: 'd', 'c', 'b', 'a' in row 1; 'h', 'g', 'f', 'e' in row 3; and 'o', ' ', ' ', ' ' in row 1023.

Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0				
1	1	0	d	c	b
2	0				
3	1	0	h	g	f
4	0				
5	0				
6	0				
7	0				
...			...		
1022	0				
1023	0				

# Cache Block 1 Tag does not match ( $0 \neq 2$ )

- 0000000000000000000010    0000000001    0100  
Tag                          Index                  Offset



# Miss, so replace block 1 with new data & tag

▪ 0000000000000000000010 0000000001 0100

	Tag	Index	Offset			
Valid	Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	2	I	k	j	i	
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...			...			
1022	0					
1023	0					

# And return word J

▪ 000000000000000010 0000000001 0100  
Tag Index Offset

Valid	Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	1	2	I	k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

# Do an example yourself. What happens?

- Chose from: Cache: Hit, Miss, Miss w. replace  
Values returned: a ,b, c, d, e, ..., k, l
- Read address **0x00000030** ?  
00000000000000000000000000000000 00000000011 0000
- Read address **0x0000001c** ?  
00000000000000000000000000000000 00000000001 1100

Index

0	0					
1	2	l	k	l	i	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					

# Answers

- **0x00000030** a hit  
Index = 3, Tag matches,  
Offset = 0, value = e
- **0x0000001c** a miss  
Index = 1, Tag mismatch, so  
replace from memory,  
Offset = **0xc**, value = d
- Since reads, values  
must = memory values  
whether or not cached:
  - 0x00000030 = e
  - 0x0000001c = d

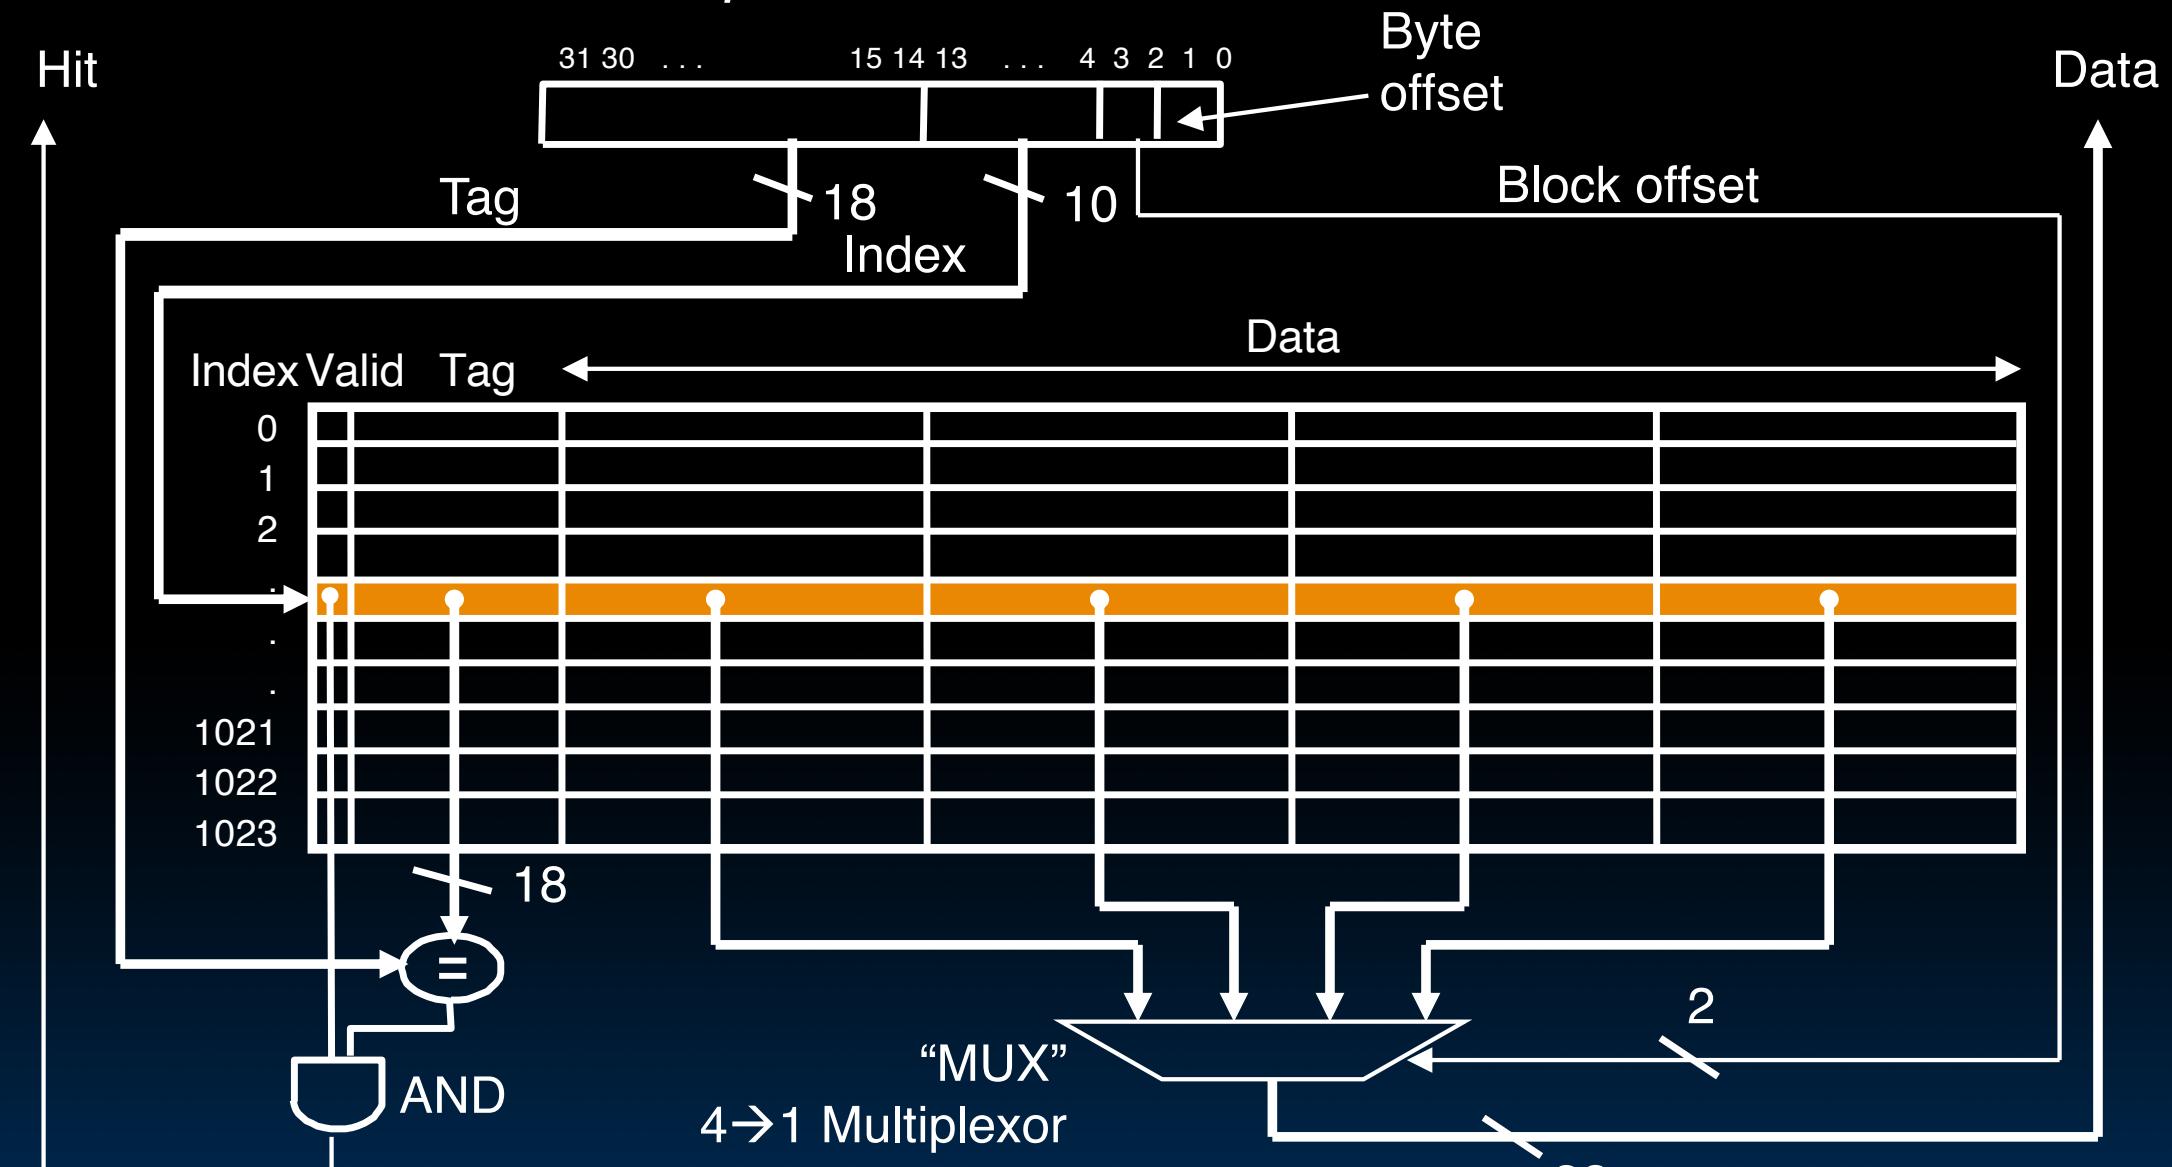
Memory Address (hex)	Value of Word
00000010	a
<u>00000014</u>	b
00000018	c
<u>0000001C</u>	d
...	...
00000030	e
<u>00000034</u>	f
00000038	g
0000003C	h
...	...
00008010	i
<u>00008014</u>	j
00008018	k
0000801C	l
...	...



Writes, Block  
Sizes, Misses

# Multiword-Block Direct-Mapped Cache

- Four words/block, cache size = 4K words



What kind of locality are we taking advantage of?

# What to do on a write hit?

- Write-through
  - Update both cache and memory
- Write-back
  - update word in cache block
  - allow memory word to be “stale”
  - add ‘dirty’ bit to block
    - memory & Cache inconsistent
    - needs to be updated when block is replaced
  - ...OS flushes cache before I/O...
- Performance trade-offs?

# Block Size Tradeoff

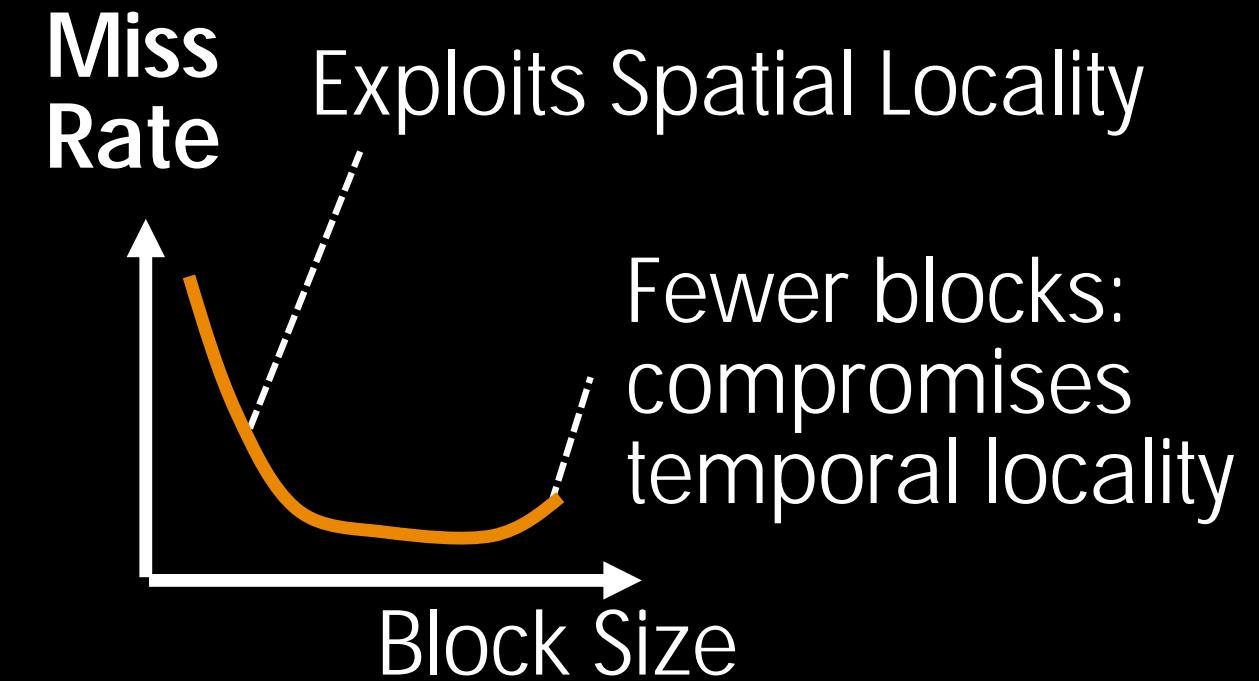
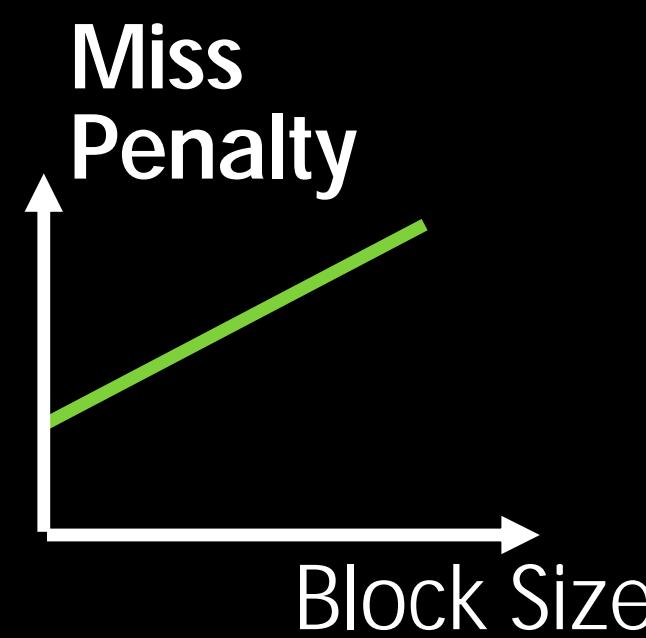
- Benefits of Larger Block Size
  - Spatial Locality: if we access a given word, we're likely to access other nearby words soon
  - Very applicable with Stored-Program Concept
  - Works well for sequential array accesses
- Drawbacks of Larger Block Size
  - Larger block size means larger miss penalty
    - on a miss, takes longer time to load a new block from next level
  - If block size is too big relative to cache size, then there are too few blocks
    - Result: miss rate goes up

# Extreme Example: One Big Block

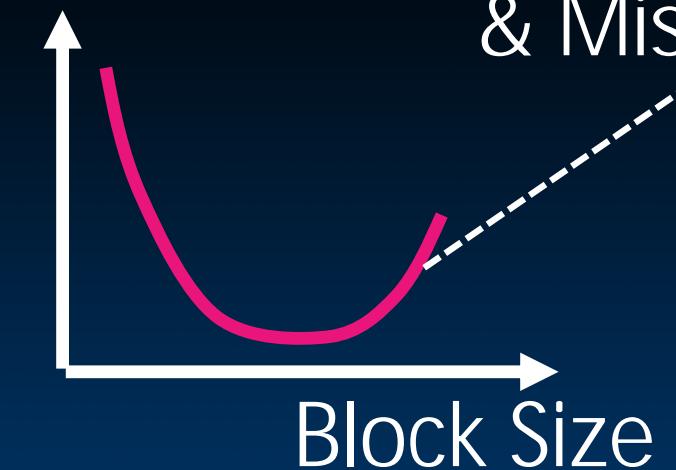


- Cache Size = 4 bytes Block Size = 4 bytes
  - Only ONE entry (row) in the cache!
- If item accessed, likely accessed again soon
  - But unlikely will be accessed again immediately!
- The next access will likely to be a miss again
  - Continually loading data into the cache but discard data (force out) before use it again
  - Nightmare for cache designer: Ping Pong Effect

# Block Size Tradeoff Conclusions



Average Access Time



# Types of Cache Misses (1/2)

- “Three Cs” Model of Misses
- 1<sup>st</sup> C: **Compulsory Misses**
  - occur when a program is first started
  - cache does not contain any of that program’s data yet, so misses are bound to occur
  - can’t be avoided easily, so won’t focus on these in this course
  - Every block of memory will have one compulsory miss (NOT only every block of the cache)

# Types of Cache Misses (2/2)

- **2<sup>nd</sup> C: Conflict Misses**
  - miss that occurs because two distinct memory addresses map to the same cache location
  - two blocks (which happen to map to the same location) can keep overwriting each other
  - big problem in direct-mapped caches
  - how do we lessen the effect of these?
- **Dealing with Conflict Misses**
  - Solution 1: Make the cache size bigger
    - Fails at some point
  - Solution 2: Multiple distinct blocks can fit in the same cache Index?



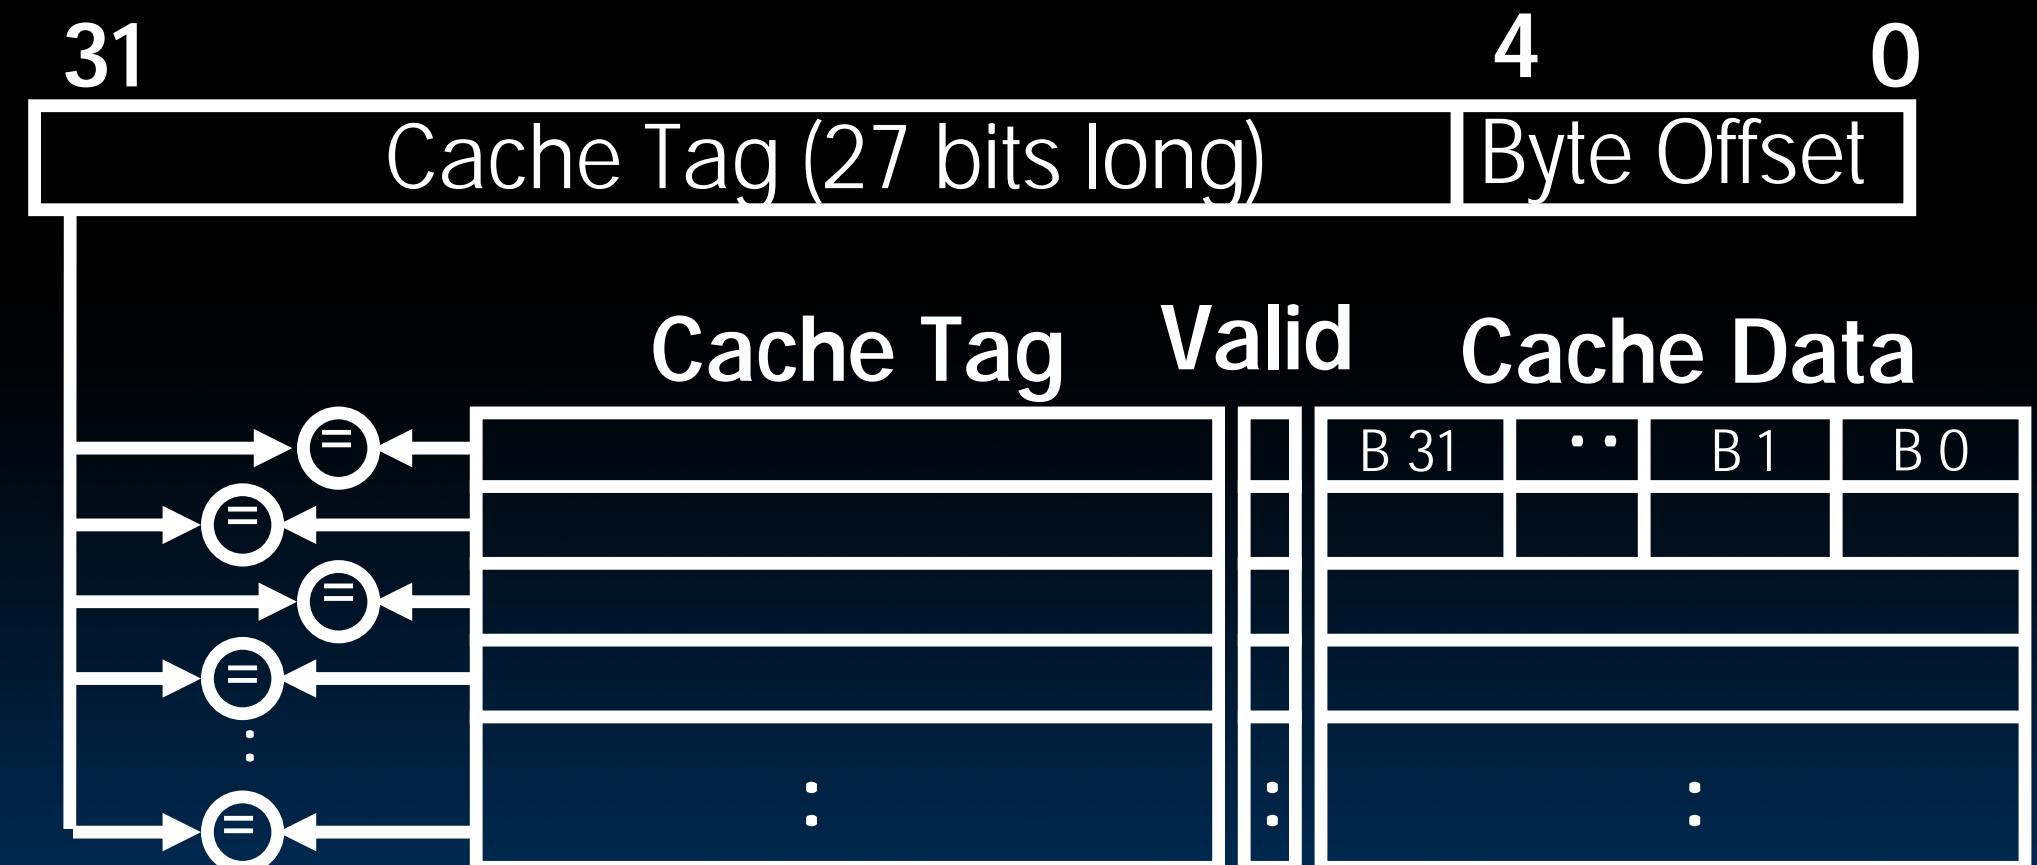
# Fully Associative Caches

# Fully Associative Cache (1/3)

- Memory address fields:
  - Tag: same as before
  - Offset: same as before
  - Index: non-existent
- What does this mean?
  - no “rows”: any block can go anywhere in the cache
  - must compare with all tags in entire cache to see if data is there

# Fully Associative Cache (2/3)

- Fully Associative Cache (e.g., 32 B block)
  - compare tags in parallel



# Fully Associative Cache (3/3)

- Benefit of Fully Assoc Cache
  - No Conflict Misses (since data can go anywhere)
- Drawbacks of Fully Assoc Cache
  - Need hardware comparator for every single entry: if we have a 64KB of data in cache with 4B entries, we need 16K comparators: infeasible

# Final Type of Cache Miss

- **3<sup>rd</sup> C: Capacity Misses**
  - miss that occurs because the cache has a limited size
  - miss that would not occur if we increase the size of the cache
  - sketchy definition, so just get the general idea
- **This is the primary type of miss for Fully Associative caches.**

# How to categorize misses

- Run an address trace against a set of caches:
  - First, consider an infinite-size, fully-associative cache. For every miss that occurs now, consider it a **compulsory miss**.
  - Next, consider a finite-sized cache (of the size you want to examine) with full-associativity. Every miss that is not in #1 is a **capacity miss**.
  - Finally, consider a finite-size cache with finite-associativity. All of the remaining misses that are not #1 or #2 are **conflict misses**.
  - (Thanks to Prof. Kubitowicz for the algorithm)

# And in Conclusion..

1. Divide into T | O bits, Go to Index = I, check valid
    1. If 0, load block, set valid and tag (COMPULSORY MISS) and use offset to return the right chunk (1,2,4-bytes)
    2. If 1, check tag
      1. If Match (HIT), use offset to return the right chunk
      2. If not (CONFLICT MISS), load block, set valid and tag, use offset to return the right chunk

